

Introduction

Motivation:

- Uncertainty estimation is essential to make reliable decisions based on the predictions of deep models, but is computationally challenging.
- It is difficult to form even a Gaussian approximation to the posterior for large deep models.
- Mean-field methods reduce the computational complexity, but yield poor estimates of the uncertainty.

Contributions:

- We propose a new **stochastic, low-rank, approximate natural-gradient (SLANG)** method for Gaussian variational inference.
- Our method estimates a “low-rank plus diagonal” covariance matrix based solely on back-propagated gradients.
- SLANG is faster and more accurate than mean-field methods, and performs comparably to state-of-the-art methods.

Natural Gradient Variational Inference

Given a deep model $p(\mathcal{D}|\theta)$ with weights θ , **Gaussian Variational Inference** computes a Gaussian approximation $q(\theta) := \mathcal{N}(\theta; \mu, \Sigma)$ to the posterior by maximizing the ELBO:

$$\mathcal{L}(\mu, \Sigma) = \mathbb{E}_q \left[\underbrace{\log p(\mathcal{D}|\theta)}_{\text{Likelihood}} + \underbrace{\log \mathcal{N}(\theta | 0, \mathbf{I}/\lambda)}_{\text{Prior}} - \underbrace{\log q(\theta)}_{\text{Approximation}} \right],$$

Gradient-based methods optimize the ELBO using the stochastic gradient updates (t is the iteration, γ_t is the learning rate)

$$\mu_{t+1} = \mu_t - \gamma_t \hat{\nabla}_{\mu} \mathcal{L}_t, \quad \Sigma_{t+1} = \Sigma_t - \gamma_t \hat{\nabla}_{\Sigma} \mathcal{L}_t.$$

Gradient descent uses Euclidean geometry and may converge slowly.

Natural Gradient methods do steepest descent in the space of realizable approximations $q(\theta)$ by optimizing on the Riemannian manifold. This is expected to converge faster and gives the update [3]

$$\mu_{t+1} = \mu_t - \beta_t \Sigma_{t+1} \hat{\nabla}_{\mu} \mathcal{L}_t \quad \Sigma_{t+1}^{-1} = (1 - \beta_t) \Sigma_t^{-1} + \beta_t \hat{\nabla}_{\Sigma} \mathcal{L}_t.$$

Both methods require storing the covariance Σ_t , which is infeasible for large models. We build upon Variational Online Gauss-Newton [4], which can be modified to learn a low-rank approximation.

Variational Online Gauss-Newton approximates the Hessian with the empirical Fisher Information matrix $\hat{\mathbf{G}}(\theta_t)$. This gives

$$\mu_{t+1} = \mu_t - \beta_t \Sigma_{t+1} [\hat{\mathbf{g}}(\theta_t) + \lambda \mu_t] \\ \Sigma_{t+1}^{-1} = (1 - \beta_t) \Sigma_t^{-1} + \beta_t [\hat{\mathbf{G}}(\theta_t) + \lambda \mathbf{I}],$$

where $\hat{\mathbf{g}}(\theta_t)$ is the gradient and

$$\hat{\mathbf{G}}(\theta_t) = \frac{1}{M} \sum_{i=1}^M g_i(\theta_t) g_i(\theta_t)^{\top}$$

is the empirical Fisher Information matrix for $p(\mathcal{D} | \theta_t)$ computed with a minibatch of size M and individual gradients $g_i(\theta_t)$.

SLANG

We approximate the covariance with a “low-rank plus diagonal” matrix

$$\Sigma_t^{-1} \approx \hat{\Sigma}_t^{-1} := \mathbf{U}_t \mathbf{U}_t^{\top} + \mathbf{D}_t,$$

where \mathbf{U}_t is a $D \times L$ matrix and \mathbf{D}_t is diagonal. The cost of storing and inverting this matrix is linear in D which is reasonable when $L \ll D$. The approximate natural gradient update for $\hat{\Sigma}_t^{-1}$ is

$$\hat{\Sigma}_{t+1}^{-1} := \mathbf{U}_{t+1} \mathbf{U}_{t+1}^{\top} + \mathbf{D}_{t+1} \approx (1 - \beta_t) \hat{\Sigma}_t^{-1} + \beta_t [\hat{\mathbf{G}}(\theta_t) + \lambda \mathbf{I}]$$

This update may increase the rank of \mathbf{U}_{t+1} , so we project the matrix onto a L -dimensional subspace using an eigenvalue decomposition:

$$(1 - \beta_t) \hat{\Sigma}_t^{-1} + \beta_t [\hat{\mathbf{G}}(\theta_t) + \lambda \mathbf{I}] = \underbrace{(1 - \beta_t) \mathbf{U}_t \mathbf{U}_t^{\top} + \beta_t \hat{\mathbf{G}}(\theta_t)}_{\text{Rank at most } L + M} + \underbrace{(1 - \beta_t) \mathbf{D}_t + \beta_t \lambda \mathbf{I}}_{\text{Diagonal component}} \\ \approx \underbrace{\mathbf{Q}_{1:L} \mathbf{\Lambda}_{1:L} \mathbf{Q}_{1:L}^{\top}}_{\text{Rank } L \text{ eigendecomposition}} + \underbrace{(1 - \beta_t) \mathbf{D}_t + \beta_t \lambda \mathbf{I}}_{\text{Diagonal component}}.$$

$$\begin{array}{c} \text{Fast eig} \\ (1 - \beta) \mathbf{U}_t \mathbf{U}_t^{\top} + \beta \mathbf{G}(\theta_t) = \mathbf{Q}_{1:L} \mathbf{\Lambda}_{1:L} \mathbf{Q}_{1:L}^{\top} + (1 - \beta) \mathbf{D}_t + \beta \lambda \mathbf{I} \approx \mathbf{U}_{t+1} \mathbf{U}_{t+1}^{\top} + \mathbf{D}_{t+1} \end{array}$$

The diagonal information lost in this projection is equal to

$$\Delta_D = \text{diag} \left[(1 - \beta) \mathbf{U}_t \mathbf{U}_t^{\top} + \beta_t \hat{\mathbf{G}}(\theta_t) - \mathbf{U}_{t+1} \mathbf{U}_{t+1}^{\top} \right].$$

We add this to \mathbf{D}_t as a diagonal correction. The final SLANG update is

$$\text{SLANG: } \mathbf{U}_{t+1} = \mathbf{Q}_{1:L} \mathbf{\Lambda}_{1:L}^{1/2} \\ \mathbf{D}_{t+1} = (1 - \beta) \mathbf{D}_t + \beta_t \lambda \mathbf{I} + \Delta_D. \\ \mu_{t+1} = \mu_t - \alpha_t \left[\mathbf{U}_{t+1} \mathbf{U}_{t+1}^{\top} + \mathbf{D}_{t+1} \right]^{-1} [\hat{\mathbf{g}}(\theta_t) + \lambda \mu_t].$$

The Algorithm

Pseudo-code for SLANG is shown in Algorithm 1. α, β are learning rates, D is denoted with a vector \mathbf{d} and \mathbf{u}_j and \mathbf{v}_j are the columns of \mathbf{U} and \mathbf{V} , respectively.

Algorithm 1: SLANG

Require: Data \mathcal{D} , hyperparameters $M, L, \lambda, \alpha, \beta$

- 1: Initialize $\mu, \mathbf{U}, \mathbf{d}$
- 2: $\delta \leftarrow (1 - \beta)$
- 3: **while** not converged **do**
- 4: $\theta \leftarrow \text{fast_sample}(\mu, \mathbf{U}, \mathbf{d})$
- 5: $\mathcal{M} \leftarrow \text{sample a minibatch}$
- 6: $[\mathbf{g}_1, \dots, \mathbf{g}_M] \leftarrow \text{backprop}(\mathcal{D}_{\mathcal{M}}, \theta)$
- 7: $\mathbf{V} \leftarrow \text{fast_eig}(\delta \mathbf{u}_1, \dots, \delta \mathbf{u}_L, \beta \mathbf{g}_1, \dots, \beta \mathbf{g}_M, L)$
- 8: $\Delta_d \leftarrow \sum_{i=1}^L \delta \mathbf{u}_i^2 + \sum_{i=1}^M \beta \mathbf{g}_i^2 - \sum_{i=1}^L \mathbf{v}_i^2$
- 9: $\mathbf{U} \leftarrow \mathbf{V}$
- 10: $\mathbf{d} \leftarrow \delta \mathbf{d} + \Delta_d + \lambda \mathbf{1}$
- 11: $\hat{\mathbf{g}} \leftarrow \sum_i \mathbf{g}_i + \lambda \mu$
- 12: $\Delta_{\mu} \leftarrow \text{fast_inverse}(\hat{\mathbf{g}}, \mathbf{U}, \mathbf{d})$
- 13: $\mu \leftarrow \mu - \alpha \Delta_{\mu}$
- 14: **end while**
- 15: **return** $\mu, \mathbf{U}, \mathbf{d}$

Algorithm 2: fast_inverse($\mathbf{g}, \mathbf{U}, \mathbf{d}$)

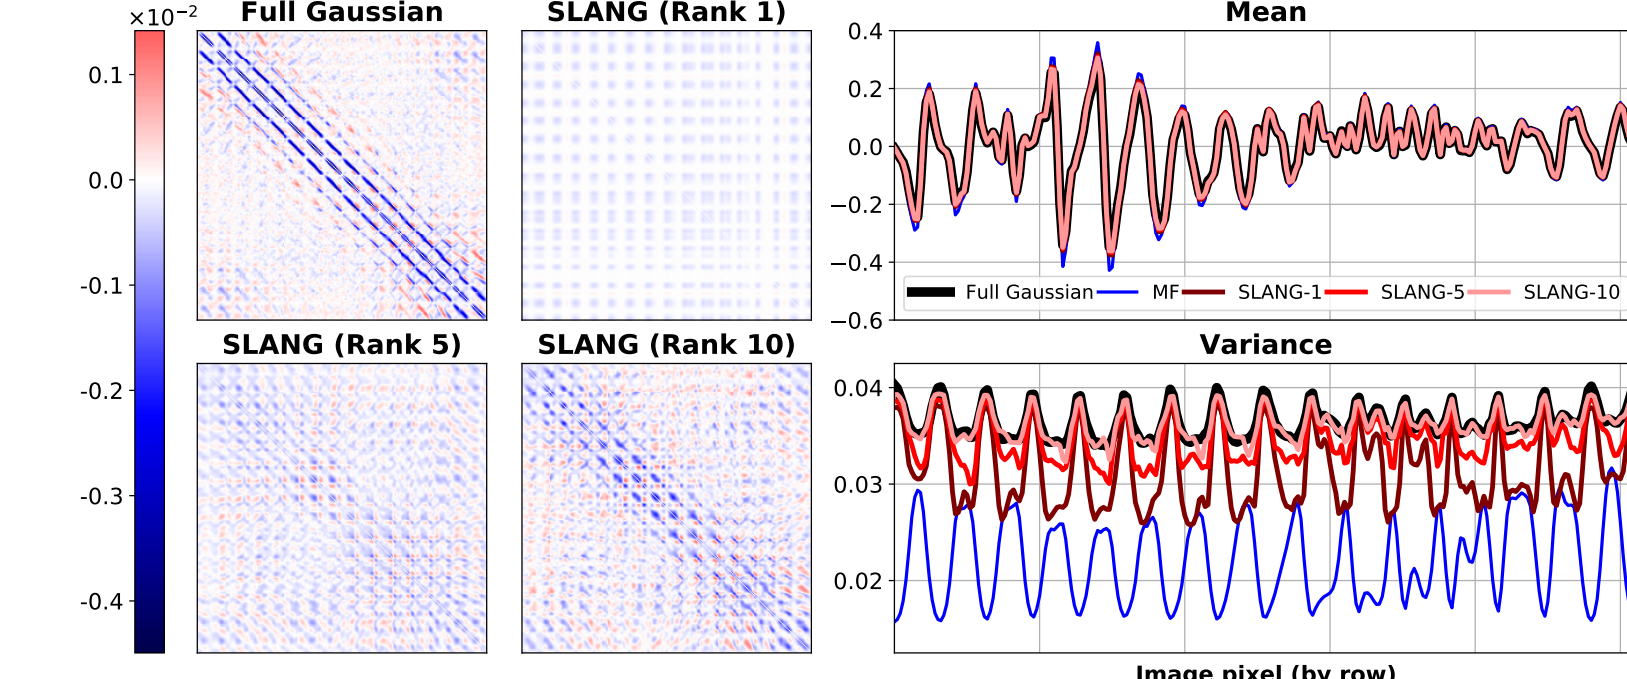
- 1: $\mathbf{A} \leftarrow (\mathbf{I}_L + \mathbf{U}^{\top} \mathbf{d}^{-1} \mathbf{U})^{-1}$
- 2: $\mathbf{y} \leftarrow \mathbf{d}^{-1} \mathbf{g} - \mathbf{d}^{-1} \mathbf{U} \mathbf{A} \mathbf{U}^{\top} \mathbf{d}^{-1} \mathbf{g}$
- 3: **return** \mathbf{y}

Algorithm 3: fast_sample($\mu, \mathbf{U}, \mathbf{d}$)

- 1: $\epsilon \sim \mathcal{N}(0, \mathbf{I}_D)$
- 2: $\mathbf{V} \leftarrow \mathbf{d}^{-1/2} \odot \mathbf{U}$
- 3: $\mathbf{A} \leftarrow \text{Cholesky}(\mathbf{V}^{\top} \mathbf{V})$
- 4: $\mathbf{B} \leftarrow \text{Cholesky}(\mathbf{I}_L + \mathbf{V}^{\top} \mathbf{V})$
- 5: $\mathbf{C} \leftarrow \mathbf{A}^{-\top} (\mathbf{B} - \mathbf{I}_L) \mathbf{A}^{-1}$
- 6: $\mathbf{K} \leftarrow (\mathbf{C} + \mathbf{V}^{\top} \mathbf{V})^{-1}$
- 7: $\mathbf{y} \leftarrow \mathbf{d}^{-1/2} \epsilon - \mathbf{V} \mathbf{K} \mathbf{V}^{\top} \epsilon$
- 8: **return** $\mu + \mathbf{y}$

Results

Covariance Structure for Logistic Regression on USPS



► SLANG doesn't underestimate variance like mean-field methods.


Logistic Regression Results

- SLANG performs similarly to full-Gaussian methods at test time.

| Dataset | Metrics | Mean-Field Methods | | | SLANG | | | Full Gaussian | | |
|------------|----------------------|--------------------|-------|-------|-------|-------|--------------|---------------|-------|-------|
| | | EF | Hess. | Exact | L = 1 | L = 5 | L = 10 | EF | Hess. | Exact |
| Australian | NLL | 0.348 | 0.347 | 0.341 | 0.342 | 0.339 | 0.338 | 0.340 | 0.339 | 0.338 |
| | KL ($\times 10^4$) | 2.240 | 2.030 | 0.195 | 0.033 | 0.008 | 0.002 | 0.000 | 0.000 | 0.000 |
| a1a | NLL | 0.339 | 0.339 | 0.339 | 0.339 | 0.339 | 0.339 | 0.339 | 0.339 | 0.339 |
| | KL ($\times 10^2$) | 2.590 | 2.208 | 1.295 | 0.305 | 0.173 | 0.118 | 0.014 | 0.000 | 0.000 |
| USPS | NLL | 0.139 | 0.139 | 0.138 | 0.132 | 0.132 | 0.131 | 0.131 | 0.130 | 0.130 |
| | KL ($\times 10^1$) | 7.684 | 7.188 | 7.083 | 1.492 | 0.755 | 0.448 | 0.180 | 0.001 | 0.000 |

Convergence Experiments

- SLANG converges faster than mean-field methods for logistic regression and BNNs.



UCI Regression with Bayesian Neural Networks:

- Performance on BNNs is comparable to Bayesian Dropout [2] and Bayes-by-Backprop [1].

| Dataset | Test RMSE | | | Test log-likelihood | | |
|----------|-----------------|-----------------------------------|-----------------------------------|-----------------------------------|------------------------------------|------------------------------------|
| | BBB | Dropout | SLANG | BBB | Dropout | SLANG |
| Boston | 3.43 \pm 0.20 | 2.97 \pm 0.19 | 3.21 \pm 0.19 | -2.66 \pm 0.06 | -2.46 \pm 0.06 | -2.58 \pm 0.05 |
| Concrete | 6.16 \pm 0.13 | 5.23 \pm 0.12 | 5.58 \pm 0.19 | -3.25 \pm 0.02 | -3.04 \pm 0.02 | -3.13 \pm 0.03 |
| Energy | 0.97 \pm 0.09 | 1.66 \pm 0.04 | 0.64 \pm 0.03 | -1.45 \pm 0.10 | -1.99 \pm 0.02 | -1.12 \pm 0.01 |
| Kin8nm | 0.08 \pm 0.00 | 0.10 \pm 0.00 | 0.08 \pm 0.00 | 1.07 \pm 0.00 | 0.95 \pm 0.01 | 1.06 \pm 0.00 |
| Naval | 0.00 \pm 0.00 | 0.01 \pm 0.00 | 0.00 \pm 0.00 | 4.61 \pm 0.01 | 3.80 \pm 0.01 | 4.76 \pm 0.00 |
| Power | 4.21 \pm 0.03 | 4.02 \pm 0.04 | 4.16 \pm 0.04 | -2.86 \pm 0.01 | -2.80 \pm 0.01 | -2.84 \pm 0.01 |

MNIST Classification with Bayesian Neural Networks:

| Test Error | SLANG | | | | | | |
|------------|-------|-------|-------|-------|-------|--------|--------------|
| | BBB | L = 1 | L = 2 | L = 4 | L = 8 | L = 16 | L = 32 |
| Test Error | 1.82% | 2.00% | 1.95% | 1.81% | 1.92% | 1.77% | 1.73% |

- Larger L leads to better test accuracy.

References

- [1] C. Blundell, J. Cornebise, K. Kavukcuoglu, and D. Wierstra. Weight uncertainty in neural networks. *CoRR*, abs/1505.05424, 2015.
- [2] Y. Gal and Z. Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. In *Proceedings of 33rd ICML*, pages 1050–1059, 2016.
- [3] M. E. Khan and D. Nielsen. Fast yet simple natural-gradient descent for variational inference in complex models. *CoRR*, abs/1807.04489, 2018.
- [4] M. E. Khan, D. Nielsen, V. Tangkaratt, W. Lin, Y. Gal, and A. Srivastava. Fast and scalable Bayesian deep learning by weight-perturbation in Adam. In *Proceedings of 35 ICML*, pages 2611–2620, 2018.